

Reading List

Jon Eskreis-Winkler

November 4, 2015

Fast Direct Methods for Gaussian Processes	Ambikasaran, Mackey, Greengard, Hogg, O’Neil	2015
--	---	------

Background: Hierarchical off-diagonal low-rank (HODLR) matrices have a k -level hierarchy depicted by this:




Fig. 1. The same HODLR matrix at different levels.

Motivation: Evaluating a Gaussian density amounts to dealing extensively with a matrix C which is typically dense.

$$p(\theta|x, y) \propto p(y|\theta, x)p(\theta) \propto \frac{1}{|\det(C(x))|^{\frac{1}{2}}} \exp\left(-\frac{1}{2}y^T C^{-1}(x)y\right) p(\theta)$$

The inversion and computation of determinant will take $O(n^3)$ operations. Result: When $C(x) = \sigma^2 I + K(x)$ for a covariance matrix K there may be a hierarchical factorization into a product of low-rank updates of the identity matrix. The resulting factorization permits inversion in $O(n \log^2 n)$ and determinant computation in $O(n \log n)$. Key ideas:

- Even though the formulation of the factorization requires inversion of large sub-matrices, there is a recursion within the large sub-matrices and in the end only small matrices get inverted.
- Use good pictures to describe the structure of a factorization
- Use Sylvester’s Theorem to speed up determinant computations.
- Use Sherman-Morrison-Woodbury Formula to invert small rank matrix updates to identity.

For additional reference see:

/Users/jeffreywinkler/Google Drive/15fall/Kondor/Presentations/ONeil

Approximating GPs with \mathcal{H}^2 matrices	Borm Garcke	2015
---	-------------	------

Background: Article deals with \mathcal{H} and \mathcal{H}^2 matrices.

- \mathcal{H} matrices are used to provide low rank approximations. For example, given a dense matrix $M \in \mathbb{R}^{m \times n}$, storage will cost $O(mn)$. However, if we are able to approximate with a rank k matrix $M \approx M^* = AB$ where $A \in \mathbb{R}^{m \times k}$, $B \in \mathbb{R}^{k \times n}$ then we can store an approximation M^* in $O(mk + nk) = O(k(m + n))$.
- \mathcal{H}^2 matrices replace the general low-rank structure of the blocks by a hierarchical representation closely related to the fast multipole method in order to reduce the storage complexity to $O(nk)$. Is this an improvement?

Motivation: Reconstructing $f : X \rightarrow Y$ using GP requires, in evaluation of the density $O(n^3)$ for direct methods and $O(n^2)$ for iterative methods. This makes large scale learning problems intractable. Result: Use Gaussian RBF kernel to make a matrix. Use \mathcal{H} matrices as data-sparse approximations of the kernel matrix. This would reduce storage for $M \in \mathbb{R}^{N \times N}$ matrices to $O(Nm \log N)$ where m is the rank of the local approximations (why the extra $\log N$?). m will also determine accuracy. Also matrix multiplications will take $O(Nm \log N)$ and inverse (and determinant) is almost linear.

Key ideas:

- Kernel matrices may be full rank, but there is a natural and surprising factorization. For a matrix $K \in \mathbb{R}^{N \times N}$, $K_{i,j} = k(x_i, x_j)$. Using m Lagrange polynomials $(\mathcal{L}_\nu)_{\nu=1}^m$ and m interpolation points $(\xi_\nu)_{\nu=1}^m$, we approximate the kernel function as $\tilde{k}(x_i, x_j) = \sum_{\nu=1}^m \mathcal{L}_\nu(x_i) k(\xi_\nu, x_j)$:

$$\tilde{K} = \begin{pmatrix} \mathcal{L}_1(x_1) & \cdots & \mathcal{L}_m(x_1) \\ \vdots & \ddots & \vdots \\ \mathcal{L}_1(x_N) & \cdots & \mathcal{L}_m(x_N) \end{pmatrix} \begin{pmatrix} k(\xi_1, x_1) & \cdots & k(\xi_1, x_N) \\ \vdots & \ddots & \vdots \\ k(\xi_m, x_1) & \cdots & k(\xi_m, x_N) \end{pmatrix}$$

When k is a sufficiently smooth kernel, m can be quite small.

- This Lagrange-based factorization can be tough when k is not globally smooth - only locally smooth. This would lead us to want to define sub-matrices to factorize in this way. The new task is to identify the sub-matrices to factor. The search is conducted by using binary space partitioning. The binary partitioning begins to be intractable for large dimensional data because there are so many border clusters. Barring this issue, finding the factorization has a complexity of $O(Nm \log N)$.

- For \mathcal{H}^2 approximation, the factorization makes a stronger assumption, that both variables are interpolated. Mathematically, $\tilde{k}(x, y) = \sum_{\nu=1}^m \sum_{\mu=1}^m \mathcal{L}_\nu(x) k(\xi_\nu, \xi_\mu) \mathcal{L}_\mu(z)$. This introduces a "coupling matrix" $S \in \mathbb{R}^{m \times m}$ and "cluster bases" $V, W \in \mathbb{R}^{N \times m}$ approximating $\tilde{K} = VSW^T =$

$$\begin{pmatrix} \mathcal{L}_1(x_1) & \cdots & \mathcal{L}_m(x_1) \\ \vdots & \ddots & \vdots \\ \mathcal{L}_1(x_N) & \cdots & \mathcal{L}_m(x_N) \end{pmatrix} \begin{pmatrix} k(\xi_1, \xi_1) & \cdots & k(\xi_1, \xi_m) \\ \vdots & \ddots & \vdots \\ k(\xi_m, \xi_m) & \cdots & k(\xi_m, x_N) \end{pmatrix} \begin{pmatrix} \mathcal{L}_1(x_1) & \cdots & \mathcal{L}_1(x_N) \\ \vdots & \ddots & \vdots \\ \mathcal{L}_m(x_1) & \cdots & \mathcal{L}_m(x_N) \end{pmatrix}$$

The complexity of this approach to approximating K is $O(Nm)$.

- How to address the curse of dimensionality? "Coarsening"